

25 YEARS ANNIVERSARY  
SOICT

HA NOI UNIVERSITY OF SCIENCE AND TECHNOLOGY  
SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY



HA NOI UNIVERSITY OF SCIENCE AND TECHNOLOGY  
SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

# Chapter 6

# Batch processing - part 1

## MapReduce

# Data processing: MapReduce

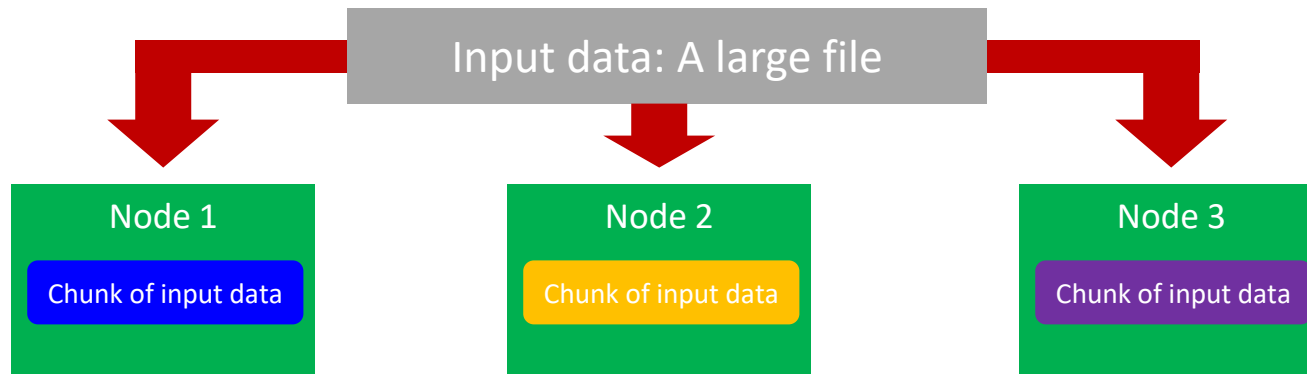
- MapReduce framework is the Hadoop default data processing engine
- MapReduce is a programming model for data processing
  - it is not a language, a style of processing data created by Google
- The beauty of MapReduce
  - Simplicity
  - Flexibility
  - Scalability

# a MR job = {Isolated Tasks}<sub>n</sub>

- MapReduce divides the workload into multiple *independent tasks* and schedule them across cluster nodes
- A work performed by each task is done *in isolation* from one another for scalability reasons
  - The communication overhead required to keep the data on the nodes synchronized at all times would prevent the model from performing reliably and efficiently at large scale


# Data Distribution

- In a MapReduce cluster, data is usually managed by a distributed file systems (e.g., HDFS)
- Move code to data and not data to code



# Keys and Values


- The programmer in MapReduce has to specify two functions, the *map function* and the *reduce function* that implement the Mapper and the Reducer in a MapReduce program
- In MapReduce data elements are always structured as  
key-value (i.e.,  $(K, V)$ ) pairs
- The map and reduce functions receive and *emit*  $(K, V)$  pairs



# Partitions

- A different subset of intermediate key space is assigned to each Reducer
- These subsets are known as *partitions*

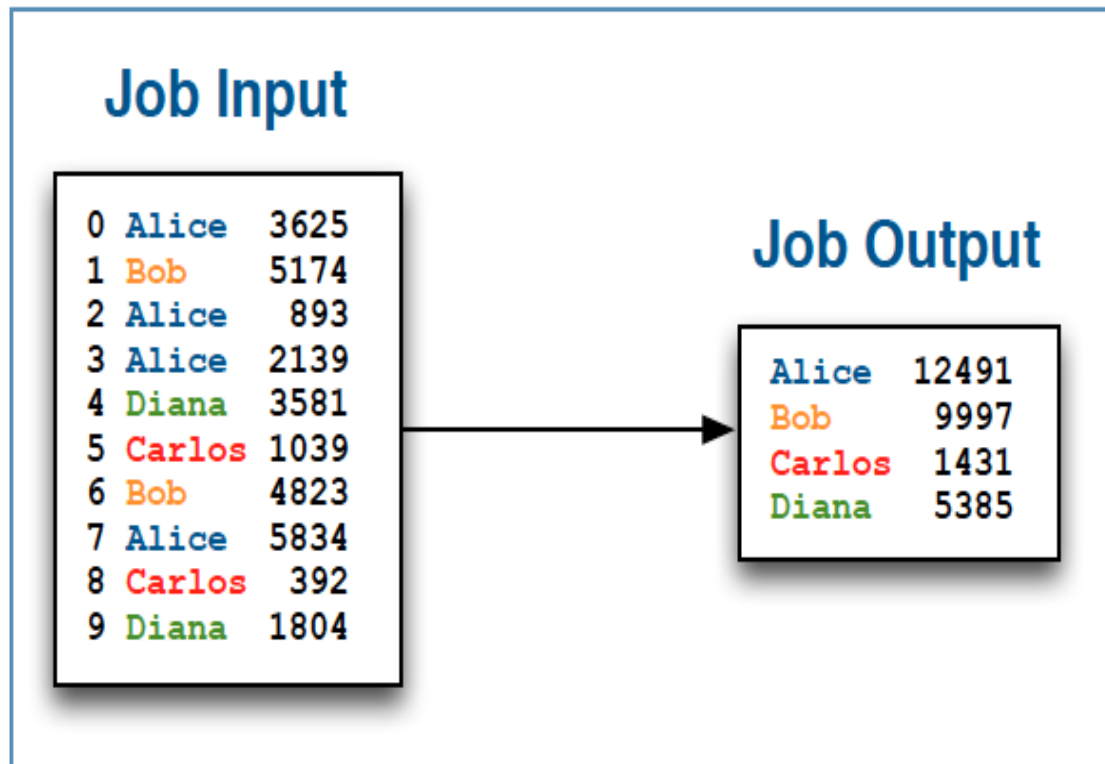
*Different colors represent different keys (potentially) from different Mappers*



*Partitions are the input to Reducers*

# MapReduce example

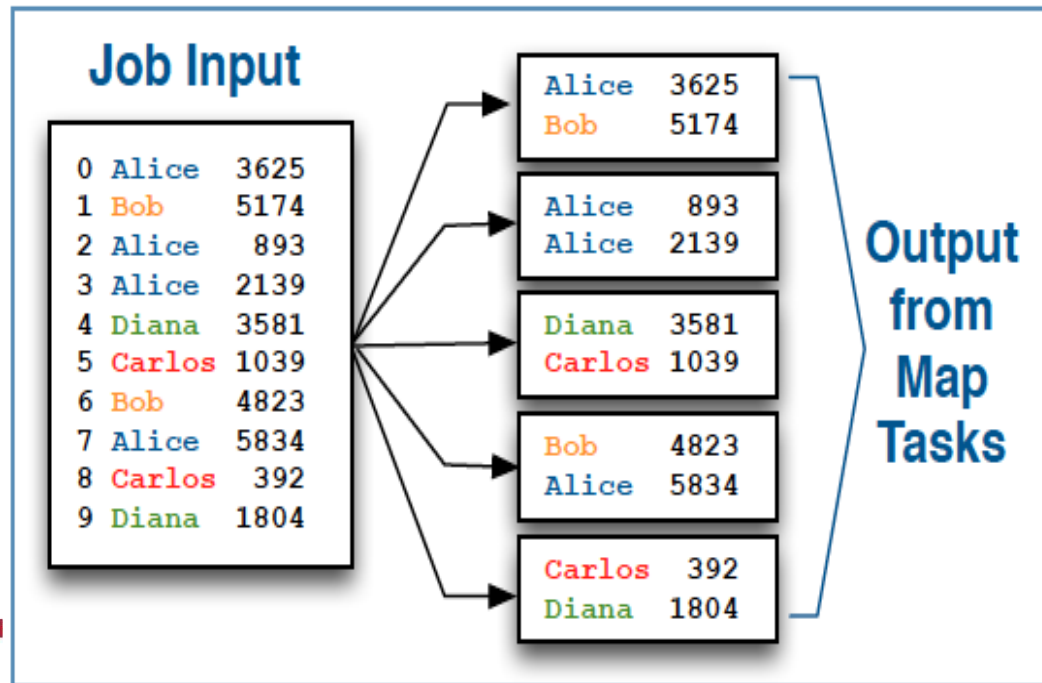
- Input: text file containing order ID, employee name, and sale amount
- Output: sum of all sales per employee





# Map phase


- Hadoop splits job into many individual map tasks
  - Number of map tasks is determined by the amount of input data
  - Each map task receives a portion of the overall job input to process
  - Mappers process one input record at a time
  - For each input record, they emit zero or more records as output
- In this case, the map task simply parses the input record
  - And then emits the name and price fields for each as output



Map phase

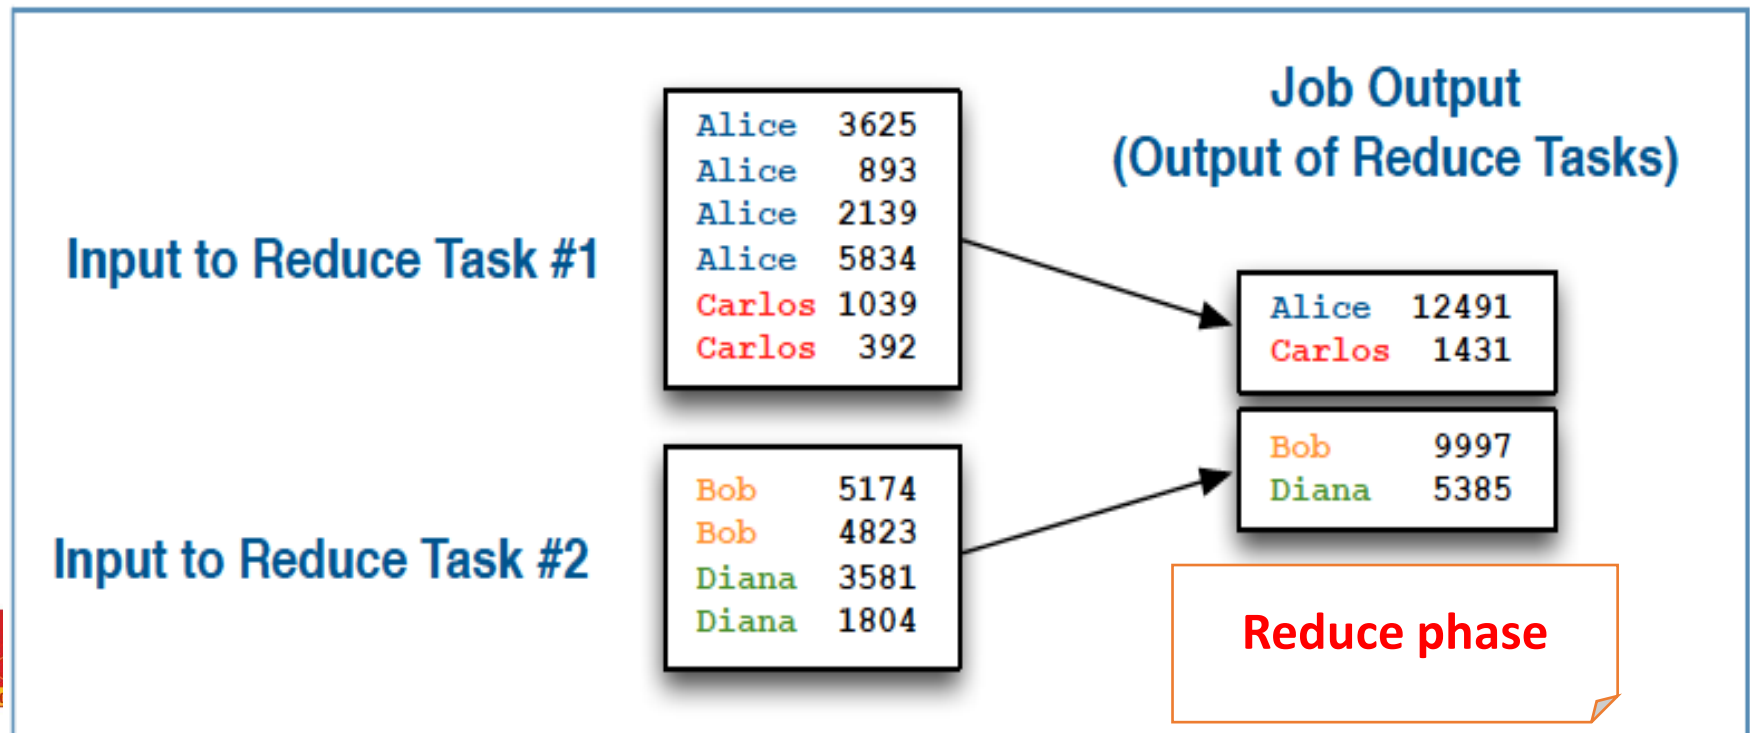
# Shuffle & sort

- Hadoop automatically sorts and merges output from all map tasks
  - This intermediate process is known as the **shuffle and sort**
  - The result is supplied to reduce tasks

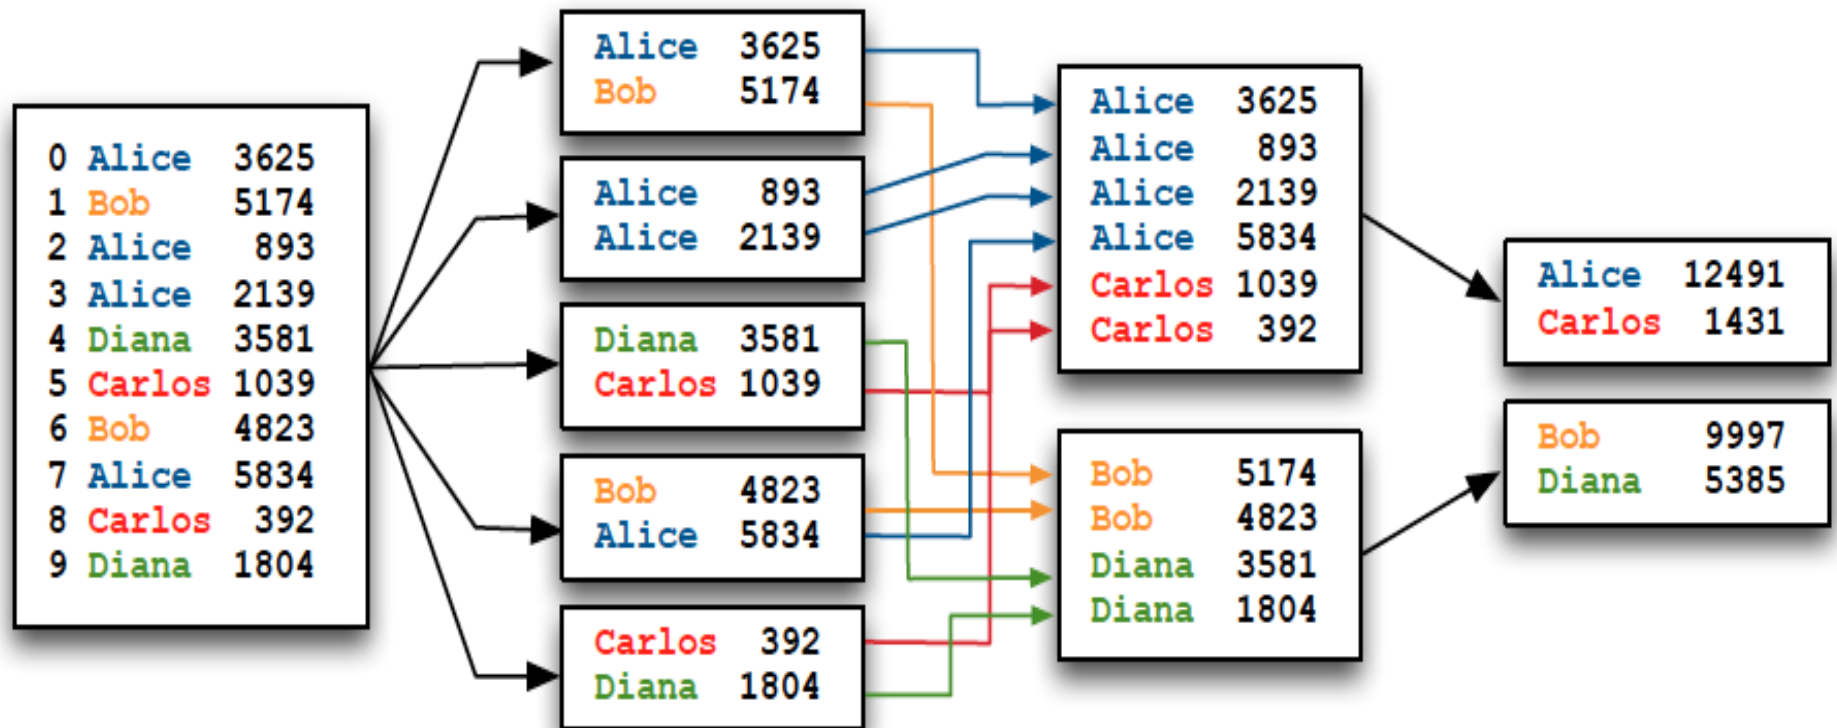


# Reduce phase

- Reducer input comes from the shuffle and sort process
  - As with map, the reduce function receives one record at a time
  - A given reducer receives all records for a given key
  - For each input record, reduce can emit zero or more output records
- Our reduce function simply sums total per person
  - And emits employee name (key) and total (value) as output

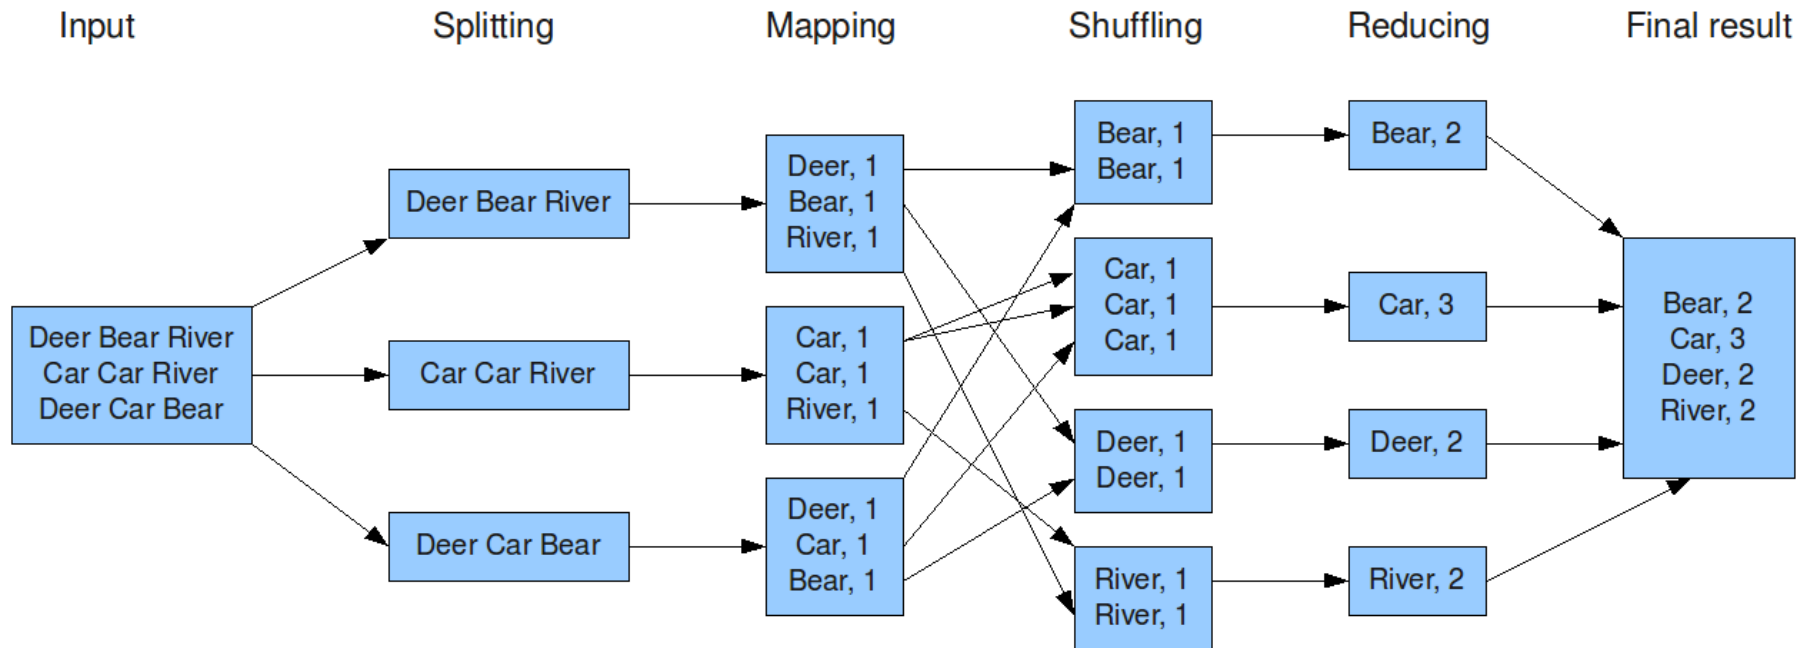


# Data flow for the entire MapReduce job

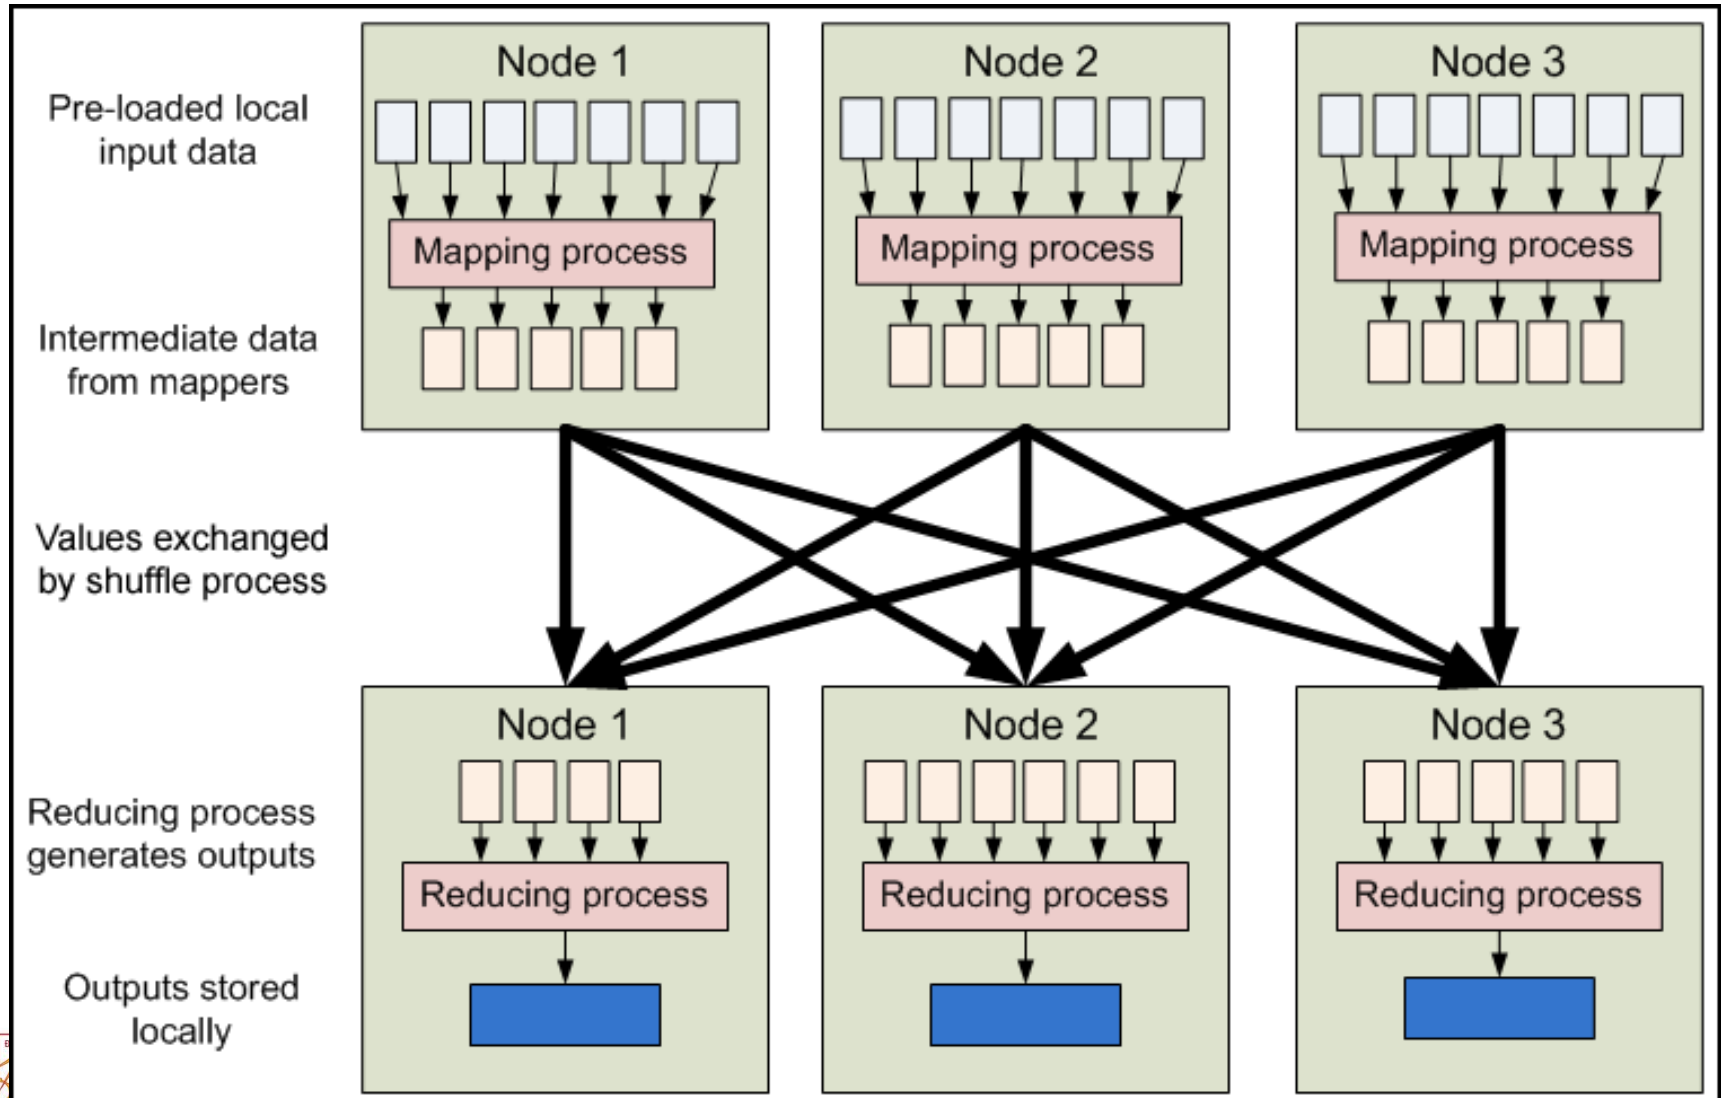


# Word Count Dataflow

The overall MapReduce word count process



# MapReduce - Dataflow



# Example: Word Count (1)

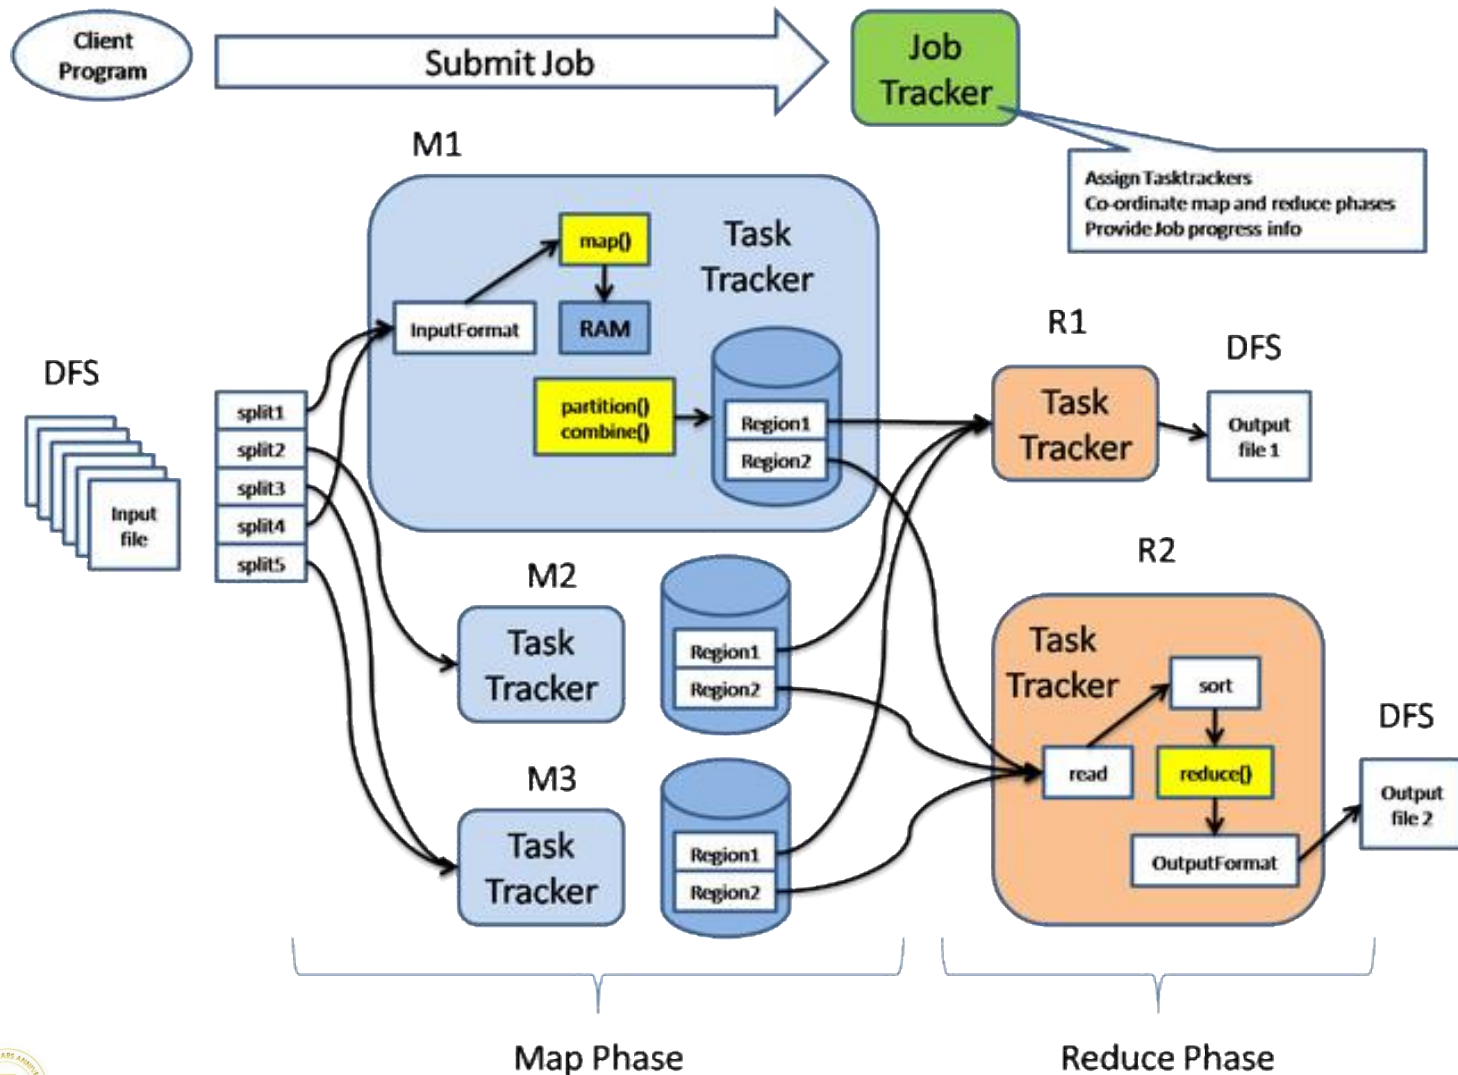
```
9 import org.apache.hadoop.mapreduce.Job;
10 import org.apache.hadoop.mapreduce.Mapper;
11 import org.apache.hadoop.mapreduce.Reducer;
12 import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
13 import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
14 import org.apache.hadoop.util.GenericOptionsParser;
15
16
17
18
19 public class WordCount {
20     public static void main(String [] args) throws Exception
21     {
22         Configuration c=new Configuration();
23         String[] files=new GenericOptionsParser(c,args).getRemainingArgs();
24         Path input=new Path(files[0]);
25         Path output=new Path(files[1]);
26         Job j=new Job(c,"wordcount");
27         j.setJarByClass(WordCount.class);
28         j.setMapperClass(MapForWordCount.class);
29         j.setReducerClass(ReduceForWordCount.class);
30         j.setOutputKeyClass(Text.class);
31         j.setOutputValueClass(IntWritable.class);
32         FileInputFormat.addInputPath(j, input);
33         FileOutputFormat.setOutputPath(j, output);
34         System.exit(j.waitForCompletion(true)?0:1);
35     }
```

# Example: Word Count (2)

```
36 public static class MapForWordCount extends Mapper<LongWritable, Text, Text, IntWritable>{
37     public void map(LongWritable key, Text value, Context con) throws IOException, InterruptedException
38     {
39         String line = value.toString();
40         String[] words=line.split(" ");
41         for(String word: words )
42         {
43             Text outputKey = new Text(word.toUpperCase().trim());
44             IntWritable outputValue = new IntWritable(1);
45             con.write(outputKey, outputValue);
46         }
47     }
48 }
49
50 public static class ReduceForWordCount extends Reducer<Text, IntWritable, Text, IntWritable>
51 {
52     public void reduce(Text word, Iterable<IntWritable> values, Context con) throws IOException, InterruptedException
53     {
54         int sum = 0;
55         for(IntWritable value : values)
56         {
57             sum += value.get();
58         }
59         con.write(word, new IntWritable(sum));
60     }
```



# Map reduce life cycle



# MapReduce algorithms

(C) <https://courses.cs.washington.edu/courses/cse490h/08au/lectures.htm>

# Algorithms for MapReduce

- Sorting
- Searching
- TF-IDF
- BFS
- PageRank
- More advanced algorithms


# Sort algorithm

- Used as a test of Hadoop's raw speed
- Essentially “IO drag race”
- Input
  - A set of files, one value per line
  - Mapper key is file name, line number
  - Mapper value is the contents of the line

# Idea

- Takes advantage of reducer properties: (key, value) pairs are processed in order by key; reducers are themselves ordered
- Mapper: Identity function for value  
 $(k, v) \rightarrow (v, \_)$
- Reducer: Identity function  $(k', \_) \rightarrow (k', \text{""})$

# Idea (2)



- (key, value) pairs from mappers are sent to a particular reducer based on  $\text{hash}(\text{key})$
- Must pick the hash function for your data such that  $k1 < k2 \Rightarrow \text{hash}(k1) < \text{hash}(k2)$

# Search algorithm

- Input
  - A set of files containing lines of text
  - A search pattern to find
- Mapper key is file name, line number
- Mapper value is the contents of the line
- Search pattern sent as special parameter

# Search algorithm

- Mapper
  - Given (filename, some text) and “pattern”, if “text” matches “pattern” output (filename, \_)
- Reducer
  - Identity function



# Optimization

- Once a file is found to be interesting, we only need to mark it that way once
- Use *Combiner* function to fold redundant (filename, \_) pairs into a single one
  - Reduces network I/O

# TF-IDF algorithm

- Term Frequency – Inverse Document Frequency
  - Relevant to text processing
  - Common web analysis algorithm

$$\text{tf}_i = \frac{n_i}{\sum_k n_k}$$

$$\text{idf}_i = \log \frac{|D|}{|\{d : t_i \in d\}|}$$

$$\text{tfidf} = \text{tf} \cdot \text{idf}$$

- $|D|$  : total number of documents in the corpus
- $|\{d : t_i \in d\}|$  number of documents where the term  $t_i$  appears (that is  $n_i \neq 0$ ).

# Obervation

- Information needed
  - Number of times term X appears in a given document
  - Number of terms in each document
  - Number of documents X appears in total number of documents

# Job 1: Word frequency in each document

- Mapper
  - Input: (docname, contents)
  - Output: ((word, docname), 1)
- Reducer
  - Sums counts for word in document
  - Outputs ((word, docname),  $n$ )
- Combiner is same as Reducer

# Job 2: Word counts for documents

- Mapper
  - Input: ((word, docname),  $n$ )
  - Output: (docname, (word,  $n$ ))
- Reducer
  - Sums frequency of individual  $n$ 's in same doc
  - Feeds original data through
  - Outputs ((word, docname), ( $n$ ,  $N$ ))
  - $N = \sum n_i$  *sums frequency*

# Job 3: Word frequency in corpus

- Mapper
  - Input:  $((\text{word}, \text{docname}), (n, N))$
  - Output:  $(\text{word}, (\text{docname}, n, N, 1))$
- Reducer
  - Number of documents where the term *word* appear  $d$
  - Outputs  $((\text{word}, \text{docname}), (n, N, d))$

# Job 4: Calculate TF-IDF

- Mapper
  - Input: ((word, docname), (n, N, d))
  - Assume D is known (or, easy MR to find it)
  - Output ((word, docname),  $TF * IDF$ )
- Reducer
  - Just the identity function

# Final thoughts on TF-IDF

- Several small jobs add up to full algorithm
- Lots of code reuse possible
  - Stock classes exist for aggregation, identity
- Jobs 3 and 4 can really be done at once in same reducer, saving a write/read cycle
- Very easy to handle medium-large scale, but must take care to ensure flat memory usage for largest scale




# Breadth-first search algorithm

- Performing computation on a graph data structure requires processing at each node
- Each node contains node-specific data as well as links (edges) to other nodes
- Computation must traverse the graph and perform the computation step
- How do we traverse a graph in MapReduce? How do we represent the graph for this?

# Breadth-first search

- Breadth-First Search is an iterated algorithm over graphs
- Frontier advances from origin by one level with each pass



# Breadth-first search & MapReduce


- Problem
  - This doesn't “fit” into MapReduce
- Solution
  - Iterated passes through MapReduce – map some nodes, result includes additional nodes which are fed into successive MapReduce passes

# Breadth-first search & MapReduce

- Problem
  - Sending the entire graph to a map task (or hundreds/thousands of map tasks) involves an enormous amount of memory
- Solution
  - Carefully consider how we represent graphs

# Graph representations

- The most straightforward representation of graphs uses references from each node to its neighbors



# Direct references

- Structure is inherent to object
- Iteration requires linked list “threaded through” graph
- Requires common view of shared memory (synchronization!)
- Not easily serializable

```
class GraphNode
{
    Object data;
    Vector<GraphNode>
        out_edges;
    GraphNode
        iter_next;
}
```

# Adjacency matrices

- Another classic graph representation.  $M[i][j] = '1'$  implies a link from node  $i$  to  $j$ .
- Naturally encapsulates iteration over nodes

0	1	2	3	4
1	0	1	0	1
2	1	0	1	1
3	0	1	0	0
4	1	0	1	0


# Adjacency matrices: Sparse representation

- Adjacency matrix for most large graphs (e.g., the web) will be overwhelmingly full of zeros.
- Each row of the graph is absurdly long
- Sparse matrices only include non-zero elements
  - 1: (3, 1), (18, 1), (200, 1)
  - 2: (6, 1), (12, 1), (80, 1), (400, 1)
  - 3: (1, 1), (14, 1)
  - ...
  - 1: 3, 18, 200
  - 2: 6, 12, 80, 400
  - 3: 1, 14
  - ...



# Finding the shortest path

- A common graph search application is finding the shortest path from a start node to one or more target nodes
- Commonly done on a single machine with Dijkstra's Algorithm
- Can we use BFS to find the shortest path via MapReduce?



This is called the single-source shortest path problem. (a.k.a. SSSP)

# Finding the shortest path: Intuition

- We can define the solution to this problem inductively:
  - $\text{DistanceTo}(\text{startNode}) = 0$
  - For all nodes  $n$  directly reachable from  $\text{startNode}$ ,  
 $\text{DistanceTo}(n) = 1$
  - For all nodes  $n$  reachable from some other set of nodes  $S$ ,
    - $\text{DistanceTo}(n) = 1 + \min(\text{DistanceTo}(m), m \in S)$

# From intuition to algorithm

- A map task receives a node  $n$  as a key, and  $(D, \text{points-to})$  as its value
  - $D$  is the distance to the node from the start
  - $\text{points-to}$  is a list of nodes reachable from  $n$
  - $\forall p \in \text{points-to}, \text{emit}(p, D+1)$
- Reduce task gathers possible distances to a given  $p$  and selects the minimum one

# Discussion

- This MapReduce task can advance the known frontier by one hop
- To perform the whole BFS, a non-MapReduce component then feeds the output of this step back into the MapReduce task for another iteration
  - Problem: Where'd the points-to list go?
  - Solution: Mapper emits (n, points-to) as well

# Blow-up and termination

- This algorithm starts from one node
- Subsequent iterations include many more nodes of the graph as frontier advances
- Does this ever terminate?
  - Yes! Eventually, routes between nodes will stop being discovered and no better distances will be found. When distance is the same, we stop
  - Mapper should emit  $(n, D)$  to ensure that “current distance” is carried into the reducer

# Adding weights

- Weighted-edge shortest path is more useful than  $\text{cost} == 1$  approach
- Simple change: points-to list in map task includes a weight 'w' for each pointed-to node
  - emit  $(p, D+wp)$  instead of  $(p, D+1)$  for each node p
  - Works for positive-weighted graph

# Comparison to Dijkstra

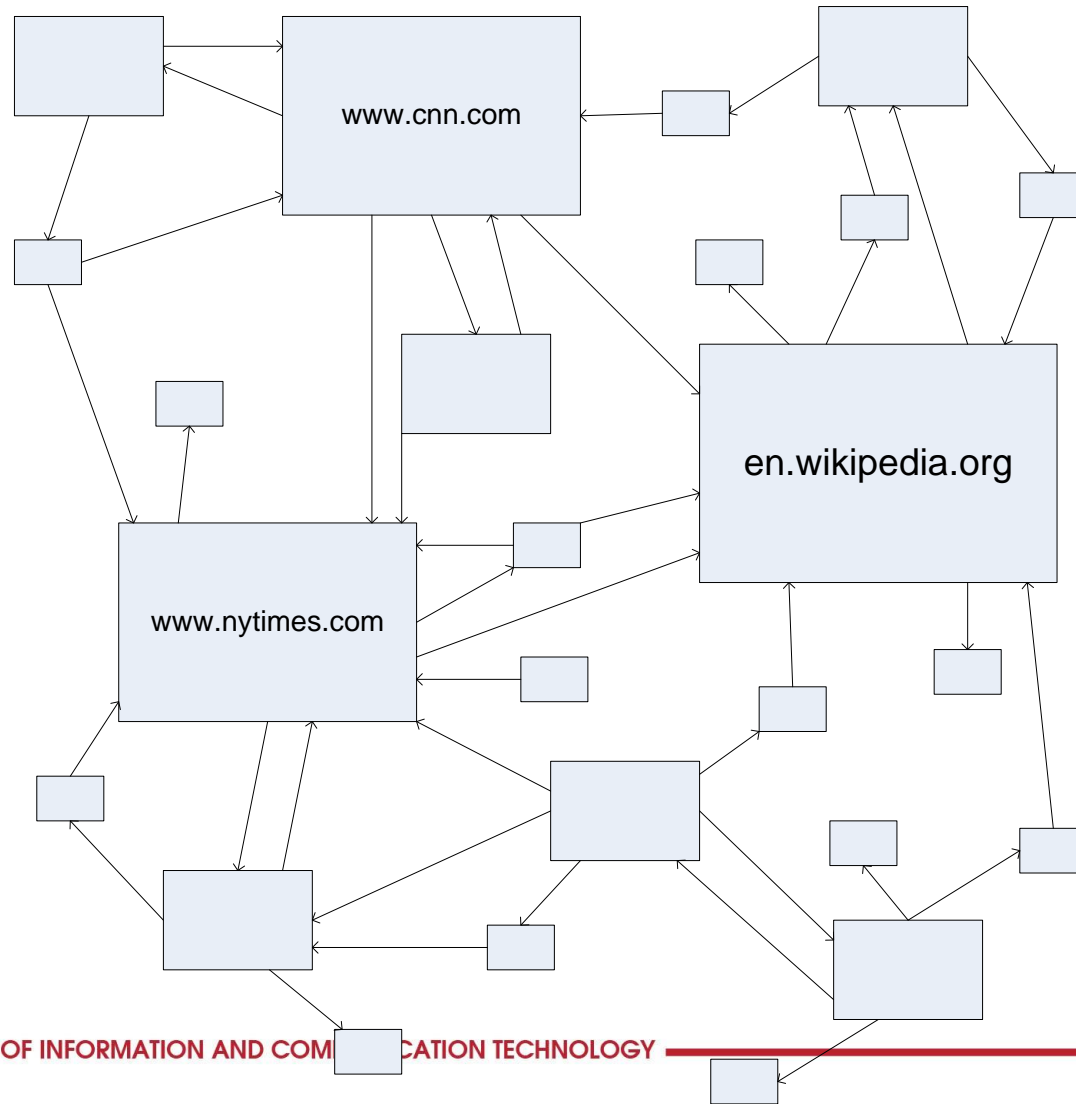
- Dijkstra's algorithm is more efficient because at any step it only pursues edges from the minimum-cost path inside the frontier
- MapReduce version explores all paths in parallel; not as efficient overall, but the architecture is more scalable
- Equivalent to Dijkstra for weight=1 case

# PageRank: Random walks over the Web

- If a user starts at a random web page and surfs by clicking links and randomly entering new URLs, what is the probability that s/he will arrive at a given page?
- The PageRank of a page captures this notion
  - More “popular” or “worthwhile” pages get a higher rank



# PageRank: Visually



# PageRank: Formula

- Given page A, and pages  $T_1$  through  $T_n$  linking to A, PageRank is defined as:
  - $PR(A) = (1-d) + d (PR(T_1)/C(T_1) + \dots + PR(T_n)/C(T_n))$
- $C(P)$  is the cardinality (out-degree) of page P
- d is the damping (“random URL”) factor

# PageRank: Intuition

- Calculation is iterative:  $PR_{i+1}$  is based on  $PR_i$
- Each page distributes its  $PR_i$  to all pages it links to. Linkees add up their awarded rank fragments to find their  $PR_{i+1}$
- $d$  is a tunable parameter (usually = 0.85) encapsulating the “random jump factor”

$$PR(A) = (1-d) + d (PR(T_1)/C(T_1) + \dots + PR(T_n)/C(T_n))$$

# PageRank: First implementation

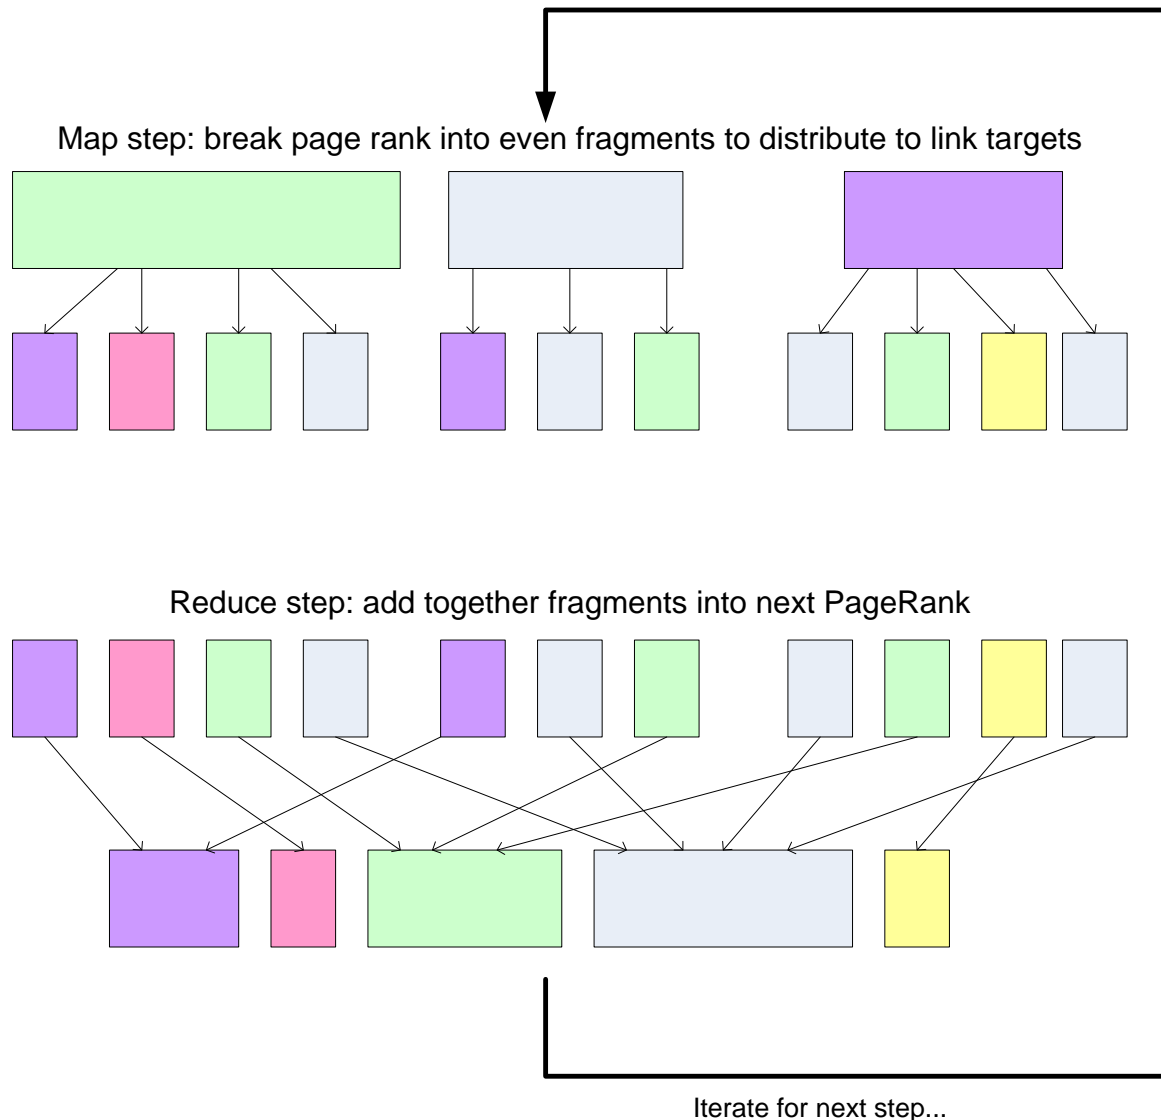
- Create two tables 'current' and 'next' holding the PageRank for each page. Seed 'current' with initial PR values
- Iterate over all pages in the graph, distributing PR from 'current' into 'next' of linkees
- $\text{current} := \text{next}; \text{next} := \text{fresh\_table}();$
- Go back to iteration step or end if converged

# Distribution of the algorithm

- Key insights allowing parallelization:
  - The 'next' table depends on 'current', but not on any other rows of 'next'
  - Individual rows of the adjacency matrix can be processed in parallel
  - Sparse matrix rows are relatively small

# Distribution of the algorithm

- Consequences of insights:
  - We can map each row of 'current' to a list of PageRank “fragments” to assign to linkees
  - These fragments can be reduced into a single PageRank value for a page by summing
  - Graph representation can be even more compact; since each element is simply 0 or 1, only transmit column numbers where it's 1



# Phase 1: Parse HTML

- Map task takes (URL, page content) pairs and maps them to (URL, (PRinit, list-of-urls))
  - PRinit is the “seed” PageRank for URL
  - list-of-urls contains all pages pointed to by URL
- Reduce task is just the identity function



# Phase 2: PageRank distribution

- Map task takes (URL, (cur\_rank, url\_list))
  - For each u in url\_list, emit (u, cur\_rank/|url\_list|)
  - Emit (URL, url\_list) to carry the points-to list along through iterations
- Reduce task gets (URL, url\_list) and many (URL, val) values
  - Sum vals and fix up with d
  - Emit (URL, (new\_rank, url\_list))

$$PR(A) = (1-d) + d (PR(T_1)/C(T_1) + \dots + PR(T_n)/C(T_n))$$

# Finishing up...

- A subsequent component determines whether convergence has been achieved (Fixed number of iterations? Comparison of key values?)
- If so, write out the PageRank lists - done!
- Otherwise, feed output of Phase 2 into another Phase 2 iteration

# Remark

- MapReduce runs the “heavy lifting” in iterated computation
- Key element in parallelization is independent PageRank computations in a given step
- Parallelization requires thinking about minimum data partitions to transmit (e.g., compact representations of graph rows)
  - Even the implementation shown today doesn't actually scale to the whole Internet; but it works for intermediate-sized graphs

# References

- Dean, Jeffrey, and Sanjay Ghemawat. "MapReduce: simplified data processing on large clusters." *Communications of the ACM* 51.1 (2008): 107-113.
- Lin, Jimmy, and Chris Dyer. "Data-intensive text processing with MapReduce." *Synthesis Lectures on Human Language Technologies* 3.1 (2010): 1-177.
- Lee, Kyong-Ha, et al. "Parallel data processing with MapReduce: a survey." *AcM SIGMoD Record* 40.4 (2012): 11-20.



25  
SOICT

VIỆN CÔNG NGHỆ THÔNG TIN VÀ TRUYỀN THÔNG  
SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Thank you  
for your  
attention!!!



[soict.hust.edu.vn/](http://soict.hust.edu.vn/)



[fb.com/groups/soict](https://fb.com/groups/soict)

