

Computer Vision

Chapter 3.2: Image transform

Nguyễn Thị Oanh
 oanhnt@soict.hust.edu.vn

Computer Vision

Chapter 3.2 Image transform


Nguyễn Thị Oanh
 oanhnt@soict.hust.edu.vn

Content

- Rappel: digital image representation
- Point Processing
- Convolution and Linear filtering
- More neighborhood operators
- Image transforms
 - Frequency domain
 - Frequencies in images
 - Fourier transform
 - Frequential Processing (frequential filters)
 - PCA (additional reading)


Frequencies in images

- Image:
 - matrix of pixels
 - a signal: a pixel ~ a sample in the image signal



Frequencies in images

- What are the (low/high) frequencies in an image?
 - Frequency = **intensity change**
 - Slow changes (homogeneous /blur regions): **low frequency**
 - fast/abrupt changes (egde, contour, noise): **high frequency**




High frequency

Low frequency

Most of energy concentrated in low frequencies

Low frequencies



High frequencies




Image spectral analysis

- An image is a visual signal
 - We can analyse the frequencies of the signal
- How?
 - we will create a **new « image »** which contains **all frequencies** of the image
 - Like a 2D frequency graphic
 - The basic tool for it is the **Fourier Transform**
- We talk about the **frequency domain**, opposing to the **spatial domain** (image)

Frequencies in a signal




FIGURE 4.1 The function at the bottom is the sum of the four functions above it. Fourier's idea in 1807 that periodic functions could be represented as a weighted sum of sines and cosines was met with skepticism.

Source: Gonzalez and Woods. *Digital Image Processing*, Prentice-Hall, 2002.

Fourier series

A bold idea (1807) - Jean

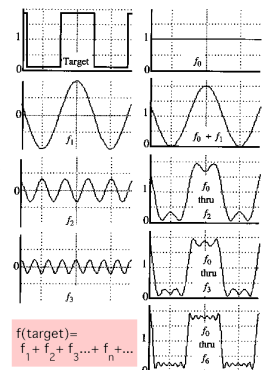
Baptiste Joseph Fourier (1768-1830):

Any univariate function can be rewritten as a **weighted sum of sines and cosines** of different frequencies.

Our building block:

$$A \sin(\omega t) + B \cos(\omega t)$$


Add enough of them to get any signal $g(t)$ you want!



Example

$$t = [0, 2], f = 1$$

$$g(t) = \sin(2\pi f t) + (1/3)\sin(2\pi(3f) t)$$




Slides: Efros


Fourier Transform

• Fourier transform is a mathematical transform that

- Decomposes functions depending on **space or time** into functions depending on **spatial or temporal frequency**




Fourier Series




We want to get this function

Fourier Series




We want to get this function

Fourier Series




We want to get this function

Fourier Series




We want to get this function


Fourier Series



Fourier Series



Fourier Series



Discrete Fourier transform

$$H_{f_j} = \frac{1}{N} \sum_k h_{t_k} e^{2\pi i f_j t_k}$$
$$h_{t_j} = \frac{1}{N} \sum_k H_{f_k} e^{-2\pi i f_k t_j}$$

where the t_k are the time corresponding to my signal in the time domain h_{t_k} , f_k are the corresponding frequency to my signal in the frequency domain, and N is the number of points of the signal data.

Fourier Transform Equation Explained:
https://www.youtube.com/watch?v=8V6Hi-kP9EE&ab_channel=IainExplainsSignals%2CSystems%2CandDigitalComms

Magnitude and phase


- Fourier transform stores the **magnitude** and **phase** at each frequency
 - Magnitude **encodes how much signal there is** at a particular frequency
 - Phase **encodes spatial** information (indirectly)
 - For mathematical convenience, this is often notated in terms of **complex number** (with real and imaginary part)

Amplitude:

$$A = \pm \sqrt{R(\omega)^2 + I(\omega)^2}$$

Phase:


$$\phi = \tan^{-1} \frac{I(\omega)}{R(\omega)}$$



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

22

Basic Fourier Transform pairs




SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Source: Václav Hlaváč - Fourier transform, in 1D and in 2D

23

Basic Fourier Transform pairs




SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Source: Václav Hlaváč - Fourier transform, in 1D and in 2D

24

Basic Fourier Transform pairs




SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Source: Václav Hlaváč - Fourier transform, in 1D and in 2D

25

Example: Music

- We think of music in terms of frequencies at different magnitudes



2D FT - discrete

Direct transform

$$F(u, v) = \frac{1}{MN} \sum_{m=0}^{M-1} \sum_{n=0}^{N-1} f(m, n) \exp \left[-2\pi i \left(\frac{mu}{M} + \frac{nv}{N} \right) \right],$$

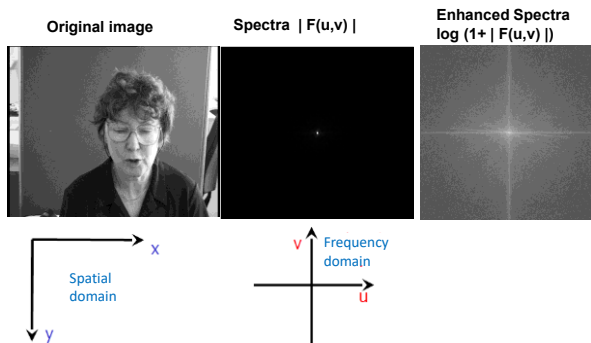
$$u = 0, 1, \dots, M-1, \quad v = 0, 1, \dots, N-1,$$

Inverse transform

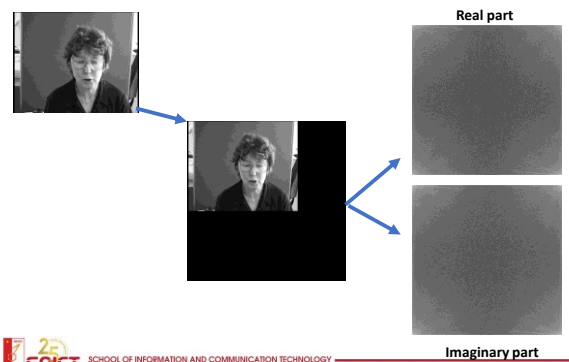
$$f(m, n) = \sum_{u=0}^{M-1} \sum_{v=0}^{N-1} F(u, v) \exp \left[2\pi i \left(\frac{mu}{M} + \frac{nv}{N} \right) \right],$$

$$m = 0, 1, \dots, M-1, \quad n = 0, 1, \dots, N-1.$$

Image Fourier transform




FFT: fast fourier transform



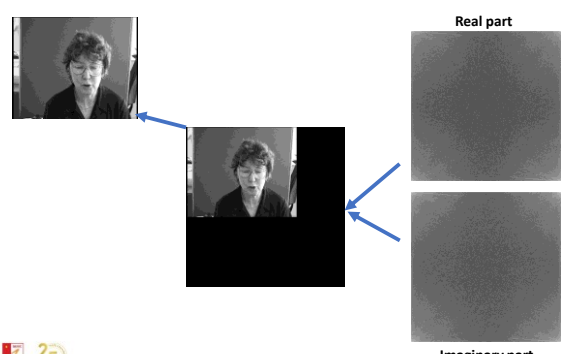
FFT

For visualization purposes we may also rearrange the quadrants of the result, so that the origin (0, 0) corresponds with the image center.




31

Inverse FFT




32

Fourier analysis in images



33


Signals can be composed



34

Fourier Bases


Teases away 'fast vs. slow' changes in the image.



This change of basis is the Fourier Transform

Hays


Fourier Bases



Discrete Fourier Transform 13

Hays

2D Fourier Transform



Slide by Steve Seltz


Fourier Transform of an image

Natural image



$f(x,y)$

Fourier decomposition
Frequency coefficients (amplitude)




$|F(\omega)|$


What does it mean to be at pixel x,y ?
What does it mean to be more or less bright in the Fourier decomposition image?

Slide by Steve Seltz

Basis reconstruction



2D Fourier transform




Percentage of image power enclosed in circles (small to large) :
90, 95, 98, 99, 99.5, 99.9


Most of energy concentrated in low frequencies

Image filtering in the frequential domain

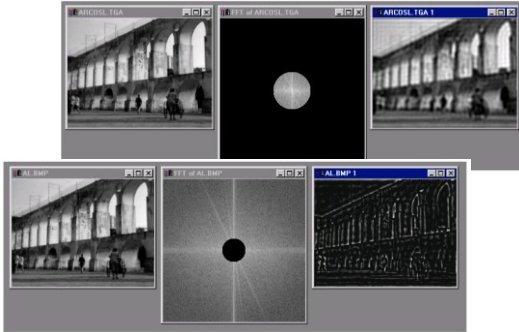
- We will be dealing only with functions (images) of finite duration so we will be interested only in Fourier Transform



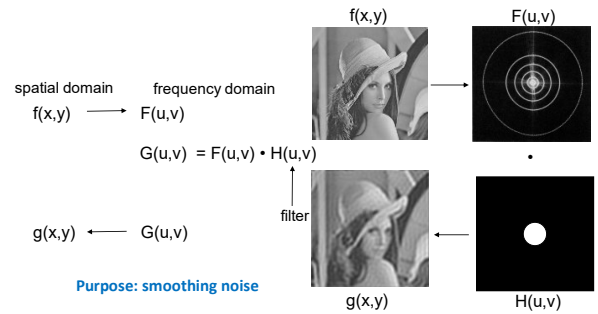
Now we can edit frequencies!



Low-pass and high-pass filtering



Low-pass filter




$H(u,v)$ - Ideal low-pass filter


$$H(u,v) = \begin{cases} 1 & D(u,v) \leq D_0 \\ 0 & D(u,v) > D_0 \end{cases}$$

$$D(u,v) = \sqrt{u^2 + v^2}$$

$$D_0 = \text{cut off frequency}$$




Blurring - Ideal low-pass filters




The blurring can be increased by increasing the size of the mask


The ringing problem



The ringing problem




$H(u,v)$ - Gaussian filter




Softer Blurring + no Ringing

Blurring - Gaussain lowpass filter




The Gaussian lowpass filter



Blurring in the Spatial Domain

Averaging = convolution with $\begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$ = point multiplication of the transform with **sinc**.


Gaussian Averaging = convolution with $\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$ = point multiplication of the transform with **a gaussian**.




High-pass filter

$H(u,v)$ - Ideal Filter

$$H(u,v) = \begin{cases} 0 & D(u,v) \leq D_0 \\ 1 & D(u,v) > D_0 \end{cases} \quad \begin{matrix} D(u,v) = \sqrt{u^2 + v^2} \\ D_0 = \text{cut off frequency} \end{matrix}$$



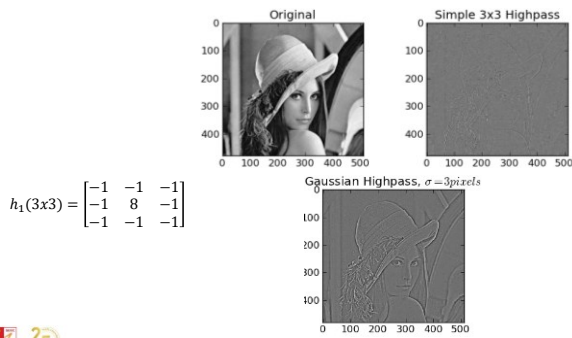
High-pass gaussian filter



$$H(u,v) = 1 - e^{-D^2(u,v)/(2D_0^2)}$$

$$D(u,v) = \sqrt{u^2 + v^2}$$

High-pass filtering



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

56

High pass filtering




FIGURE 4.11 (a) An image of size 500×500 pixels and (b) its Fourier spectrum. The superimposed circles have radii values of 5, 15, 30, 80, and 230, which enclose 92.0, 94.6, 96.4, 98.0, and 99.5% of the image power, respectively.



Source : Gonzalez and Woods. Digital Image Processing. Prentice-Hall, 2002.

57

High pass filtering




FIGURE 4.24 Results of ideal highpass filtering the image in Fig. 4.11(a) with $D_0 = 15, 30$, and 80 , respectively. Problems with ringing are quite evident in (a) and (b).




Source : Gonzalez and Woods. Digital Image Processing. Prentice-Hall, 2002.

58

Band-pass filtering

$$H(u,v) = \begin{cases} 0 & D(u,v) \leq D_0 - \frac{w}{2} \\ 1 & D_0 - \frac{w}{2} \leq D(u,v) \leq D_0 + \frac{w}{2} \\ 0 & D(u,v) > D_0 + \frac{w}{2} \end{cases}$$

$D(u,v) = \sqrt{u^2 + v^2}$
 D_0 = cut off frequency
 w = band-width




Can be obtained by multiplying the filter functions of a low-pass and of a high-pass in the frequency domain




SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

59


Removing sinus noise



Removing sinus noise



High-pass filtering + orientation



Hybrid Images



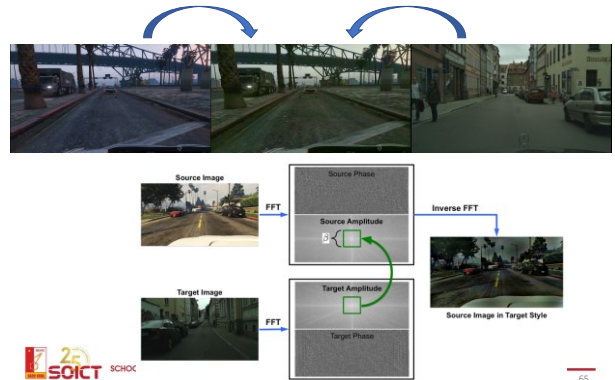


Figure 2: hybrid images are generated by superimposing two images at two different spatial scales: the low-spatial scale is obtained by filtering one image with a low-pass filter, and the high spatial scale is obtained by filtering a second image with a high-pass filter. The final hybrid image is composed by adding these two filtered images.

A. Oliva, A. Torralba, P.G. Schyns, SIGGRAPH 2006



Style transfer



Content

- Rappel: digital image representation
- Point Processing
- Convolution and Linear filtering
- More neighborhood operators
- Image transforms
 - Frequency domain
 - PCA (additional reading)
 - PCA
 - Example of using PCA for face recognition

Principle Component Analysis - PCA (Karhunen-Loeve transformation)

- **PCA** transforms the original input space into a lower dimensional space
 - By constructing dimensions that are linear combinations of the given features
- The objective: consider **independent dimensions** along which data have **largest variance** (i.e., greatest variability)

Slide by Jana Kosecka

Principal Component Analysis (cont.)

- **PCA** enables transform a number of possibly correlated variables into a smaller number of uncorrelated variables called **principal components**
- The **first principal component** accounts for as much of the **variability** in the data as possible
- Each **succeeding component** (orthogonal to the previous ones) accounts for as much of the remaining variability as possible



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Slide by Jana Kosecka

68

Principal Component Analysis (cont.)

- PCA is the most commonly used dimension reduction technique.
- Data samples

$$x_1, \dots, x_N$$

- Compute the mean

$$\bar{x} = \frac{1}{N} \sum_{i=1}^N x_i$$

- Compute the covariance matrix:

$$\Sigma_x = \frac{1}{N} \sum_{i=1}^N (x_i - \bar{x})(x_i - \bar{x})^T$$



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Slide by Jana Kosecka

69

Principal Component Analysis (cont.)

- Compute the eigenvalues λ and eigenvectors e of the matrix Σ_x
- Solve $\Sigma_x x = \lambda x$
- Order them by magnitude:

$$\lambda_1 \geq \lambda_2 \geq \dots \lambda_N.$$
- PCA reduces the dimension by keeping direction e such that $\lambda < T$.



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Slide by Jana Kosecka


70

Principal Component Analysis (cont.)

- For many datasets, most of the eigenvalues are negligible and can be discarded.

The eigenvalue λ measures the variation in the direction of corresponding eigenvector

Example:
 $\lambda_1 \neq 0, \lambda_2 = 0.$



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Slide by Jana Kosecka

71

Principal Component Analysis (cont.)

- How to get uncorrelated components which Capture most of the variance
- Project the data onto the selected eigenvectors:


$$y_i = e_i^T (x_i - \bar{x})$$
- If we consider first M eigenvectors we get new lower dimensional representation

$$[y_1, \dots, y_M]$$
- Proportion covered by first M eigenvalues

$$\frac{\sum_{i=1}^M \lambda_i}{\sum_{i=1}^N \lambda_i}$$

Slide by Jana Kosecka


Illustration of PCA



First principal component of a two-dimensional data set.

Slide by Jana Kosecka

Illustration of PCA



Second principal component of a two-dimensional data set.

Slide by Jana Kosecka

Determining the number of components

- Plot the eigenvalues
 - each eigenvalue is related to the amount of variation explained by the corresponding axis (eigenvector)
 - If the points on the graph tend to level out (show an “elbow” shape), these eigenvalues are **usually close enough to zero that they can be ignored**

Slide by Jana Kosecka

Content

- Rappel: digital image representation
- Point Processing
- Convolution and Linear filtering
- More neighborhood operators
- Image transforms
 - Frequency domain
 - PCA (additional reading)
 - PCA
 - Example of using PCA for face recognition

The space of all face images

- When viewed as **vectors of pixel values**, face images are extremely high-dimensional
 - 100x100 image = 10,000 dimensions
- However, relatively **few 10,000-dimensional vectors correspond to valid face images**
- We want to effectively **model the subspace** of face images



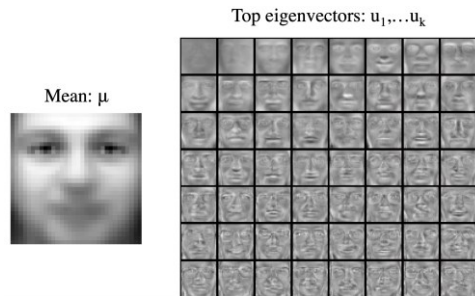
Eigenfaces: Key idea

- Assume that most face images lie on a **low-dimensional subspace determined by the first k ($k < d$) directions of maximum variance**
- Use PCA to determine the vectors or “**eigenfaces**” **u_1, \dots, u_k** that span that subspace
- Represent all face images in the dataset as **linear combinations of eigenfaces**

Eigenfaces example- Training images

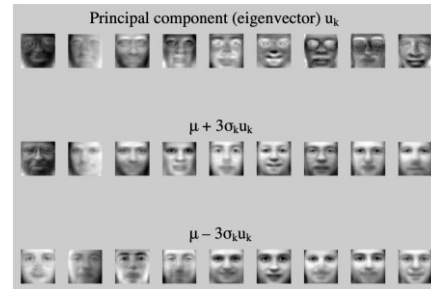


Eigenfaces example



Slide by Jana Kosecka

Eigenfaces example



Slide by Jana Kosecka

Eigenfaces examples

• Representation



$$(w_{i1}, \dots, w_{ik}) = (u_1^T(x_i - \mu), \dots, u_k^T(x_i - \mu))$$

• Reconstruction



$$\hat{x} = \mu + w_1 u_1 + w_2 u_2 + w_3 u_3 + w_4 u_4 + \dots$$

Slide by Jana Kosecka

Recognition with eigenfaces

- Process labeled training images:
- Find mean μ and covariance matrix Σ
- Find k principal components (eigenvectors of Σ) u_1, \dots, u_k
- Project each training image x_i onto subspace spanned by principal components:
 $(w_{i1}, \dots, w_{ik}) = (u_1^T(x_i - \mu), \dots, u_k^T(x_i - \mu))$
- Given novel image x :
- Project onto subspace:
 $(w_1, \dots, w_k) = (u_1^T(x - \mu), \dots, u_k^T(x - \mu))$
- Optional: check reconstruction error $x - \hat{x}$ to determine whether image is really a face
- Classify as closest training face in k -dimensional subspace

Slide by Jana Kosecka

